



AARHUS
UNIVERSITET

23 FEBRUAR 2011

NOVEL APPROACHES TO THE INDEXING OF MOVING OBJECT TRAJECTORIES

DIETER PFOSER, CHRISTIAN S. JENSEN, YANNIS THEODORIS




OVERVIEW

- › Data: Trajectories
- › Queries
- › Data structures
 - › R-tree
 - › Spatio-Temporal R-tree (STR-tree)
 - › Trajectory Bundle tree (TB-tree)
- › Query handling
- › Performance


DATA: TRAJECTORIES (1)

- > Sample position of object at time points and interpolate between samples



DATA: TRAJECTORIES (2)

- › Objects are tracked on a
- › Add time along the vertic
- › Represent in 3D:
 - › 2D spatial
 - › 1D temporal



- › The resulting 3D polyline consisting of **segments** is the **trajectory** of the moving object

IMPLICATIONS

- › Typical: Dataset → **Objects**
- › Now: Dataset → **Trajectories** → **Segments**
- › By considering time and trajectories,
we can **derive** further information:
Speed of object
Is object entering or leaving?
Was it here at the same time as another object?
Where did object go for the entire day?

QUERIES

› Coordinate-based queries

- › Point
- › Range
- › Nearest neighbour


› Trajectory-based queries

- › Topological: Enter, leave, cross, bypass
- › (Navigational: Using derived information, e.g. speed, heading...)

› Combined queries


TOPOLOGICAL QUERIES

- › Trajectories **enter**, **leave**, **cross**, **stay within**, or **bypass** a given spatiotemporal range
- › Signature: $\text{range} \times \{\text{trajectories}\} \rightarrow \{\text{trajectories}\}$




TOPOLOGICAL QUERIES

“Which taxis **left** **Tucson** **between 7 and 8 a.m. today?**”



(NAVIGATIONAL QUERIES)

- › Considering derived information in a query,
e.g., **speed** (top, average), **heading**, **traveled distance**,
covered area, etc.
- › Signature: $\text{range} \times \{\text{trajectories}\} \rightarrow \text{int} \mid \text{real} \mid \text{bool}$



COMBINED QUERIES


› Dataset → **Trajectories** → **Segments**

› *First*, selecting the *trajectory*,
and *subsequently* selecting
the *segments*

› Trajectory selection


- › trajectory identifier
- › any type of query

“What were the trajectories of the taxis that
were in **Tucson between 7 and 8 a.m. today**
until either noon or leaving Arizona?”



COMBINED QUERIES

“What were the trajectories of the taxis **until 12 p.m.**
after they **left Tucson between 7 and 8 a.m. today?**”



OVERVIEW


- › Data: Trajectories
- › Queries
- › **Data structures:**
 - › **R-tree**
 - › **Spatio-Temporal R-tree (STR-tree)**
 - › Trajectory Bundle tree (TB-tree)
- › Query handling
- › Performance

DATA STRUCTURES

› We store **segments** of **trajectories**

› R-tree

- › approximating **bounding box**
- › **orientation** of the segment, and
- › **trajectory id**

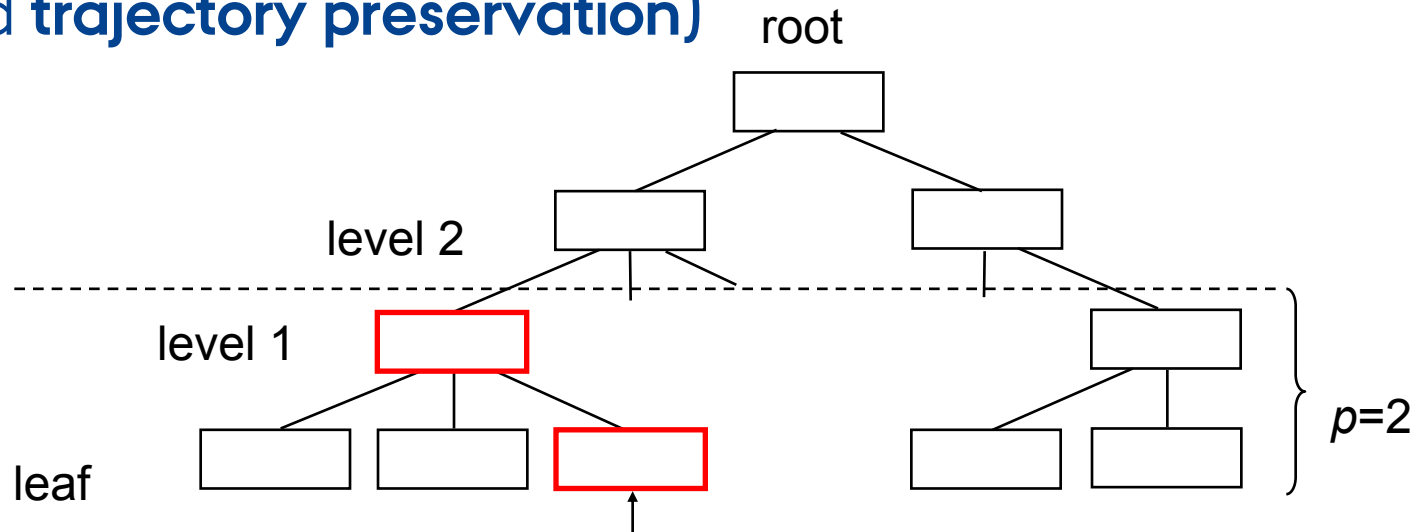


› STR-tree (SpatioTemporal R-tree)

- › **spatial discrimination**, i.e., preserve spatial proximity of segments in a leaf node, and **trajectory preservation**, i.e., segments belonging to the same trajectory
- › The tree structures are identical – the strategies for inserting and splitting differ


STR-TREE: INSERTION (1)

- › Insert new segment in leaf node containing its trajectory predecessor
- › If node full: Parameter p specifies how far back to search for a parent node to split (trade-off between **spatial packing** and **trajectory preservation**)




STR-TREE: INSERTION (2)

- › Insert the segment in the leaf node that needs least expansion, like in R-tree
- › Exclude the already traversed part of the tree



STR-TREE: NODE SPLITTING

- › 3 different node types require 3 split strategies
 - (a) quadratic split – like in R-tree
 - (b) disconnected segments in new node - *separating trajectories*
 - (c) most recent segment in new node - *continuing trajectories*



OVERVIEW

- › Data: Trajectories
- › Queries
- › Data structures:
 - › R-tree
 - › Spatio-Temporal R-tree (STR-tree)
 - › **Trajectory Bundle tree (TB-tree)**
- › Query handling
- › Performance

TB-TREE


› TB-tree (Trajectory Bundle)

› **strict trajectory preservation,**

one leaf node contains segments of only one trajectory



› **neglecting spatial discrimination**

with respect to the two spatial dimensions




TB-TREE

- > A leaf node contains a (partial) trajectory
- > Leaf nodes are additionally connected in a linked list to allow easy traversal of trajectories



TB-TREE: INSERTION

- › Try to insert segment in leaf node containing previous segment
- › Strict trajectory preservation; one leaf node, one trajectory
→ filling up nodes and no node splitting necessary
- › If full, create new with new segment as only entry




OVERVIEW

- › Data: Trajectories
- › Queries
- › Data structures:
 - › R-tree
 - › Spatio-Temporal R-tree (STR-tree)
 - › Trajectory Bundle tree (TB-tree)
- › **Query handling**
- › Performance evaluation

COMBINED QUERIES

“What were the trajectories of the taxis **until 12 p.m.**
after they **left Tucson between 7 and 8 a.m. today?**”




COMBINED QUERY HANDLING

> R- and STR-tree

- > Perform initial range query
- > Expand trajectories by looking in current leaf.
If no connecting segment: Perform new range search.

> TB-tree

- > Perform initial range query
- > Expand trajectories using the linked list (free!)



OVERVIEW

- › Data: Trajectories
- › Queries
- › Data structures:
 - › R-tree
 - › Spatio-Temporal R-tree (STR-tree)
 - › Trajectory Bundle tree (TB-tree)
- › Query handling
- › **Performance evaluation**


INDEX SIZE

- › STR and TB-tree: Nodes are full, because new segments are appended to their trajectories when possible -> small index
- › TB-tree saves a little space because we only need to store one trajectory ID pr leaf

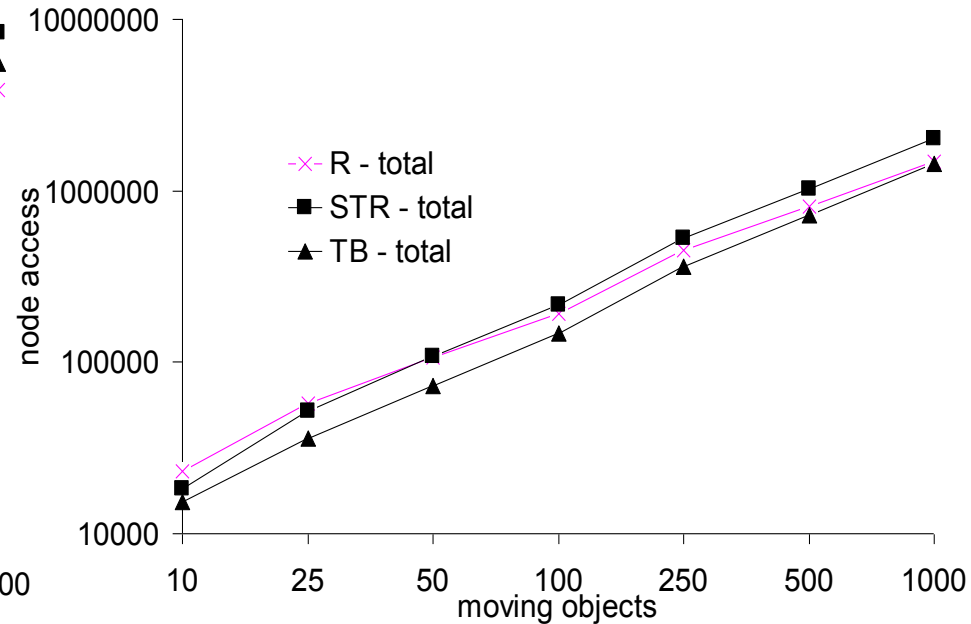
	R-tree	STR-tree	TB-tree
Index Size	~ 95 KB per object	~ 57 KB per object	~ 51 KB per object
Space Utilization	55%-60%	~100%	~100%

RANGE QUERIES

- › Increasing number of moving objects
- › less emphasis on temporal discrimination, need for spatial discrimination



Query: 10% range in each dimension



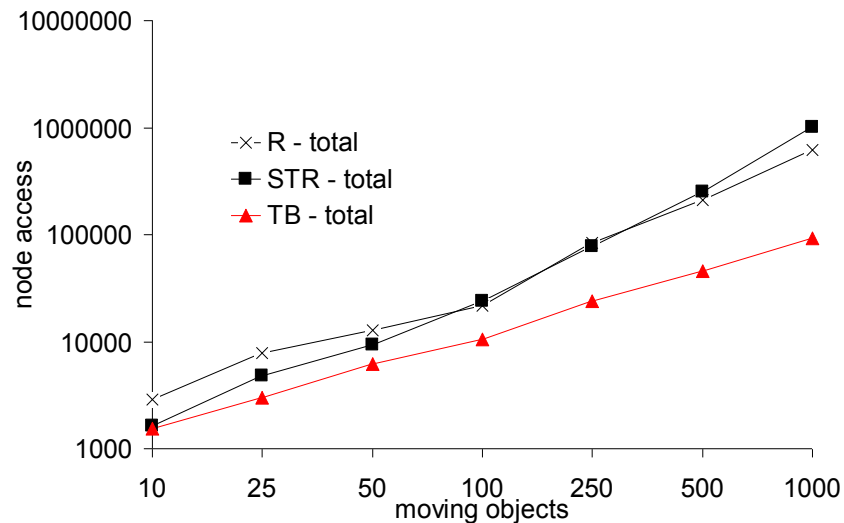
Query: 20% range in each dimension

COMBINED QUERIES


- › **TB-tree**: linked list reduces node access once the trajectory is identified
- › **STR vs. R**:

With low number of objects, the more tightly packed STR-tree is better.

With high number, better spatial discrimination of R-tree wins.



1% inner range, 10% outer



1% inner range, 20% outer

SUMMARY

- › Definition of **trajectory data** and **queries**
- › Proposal of the **STR-tree** and the **TB-tree**
- › Performance
 - › the TB-tree performs generally better than the STR-tree
 - › the TB-tree is competitive to the R-tree for coordinate-based queries
 - › the TB-tree outperforms the other methods for combined queries